

Shopify

May 22, 2022

1 Question 1

1.1 Introduction

Like any other python project, we begin by importing some important libraries and modules:

```
[2]: import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
```

We begin by reading the data:

```
[3]: d1 = pd.read_csv("2019-Winter-Data-Science-Intern-Challenge-Data-Set-Sheet1.
↪csv")
d1.head(2)
```

```
[3]:   order_id  shop_id  user_id  order_amount  total_items  payment_method \
0         1        53      746           224           2           cash
1         2        92      925           90           1           cash

      created_at
0  2017-03-13 12:36:56
1  2017-03-03 17:38:52
```

1.2 Exploratory Data Analysis:

It is a good practice to conduct a basic exploratory analysis to get a better sense of data.

```
[4]: d1.size # there are 35000 data entries
```

```
[4]: 35000
```

```
[5]: d1.describe()# gives general statistics such as count, mean, standard
↪devaiation, minimum, maximum value etc.
```

```
[5]:
```

	order_id	shop_id	user_id	order_amount	total_items
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	50.078800	849.092400	3145.128000	8.78720
std	1443.520003	29.006118	87.798982	41282.539349	116.32032

min	1.000000	1.000000	607.000000	90.000000	1.000000
25%	1250.750000	24.000000	775.000000	163.000000	1.000000
50%	2500.500000	50.000000	849.000000	284.000000	2.000000
75%	3750.250000	75.000000	925.000000	390.000000	3.000000
max	5000.000000	100.000000	999.000000	704000.000000	2000.000000

order_amount seems to have the an unusually high standard deviation. the range is wide from 90 units to 704000 units. We could use this information later.

```
[6]: d1.count()
```

```
[6]: order_id      5000
shop_id      5000
user_id      5000
order_amount  5000
total_items   5000
payment_method 5000
created_at    5000
dtype: int64
```

```
[8]: d1.isna().sum() # No null Values.
```

```
[8]: order_id      0
shop_id      0
user_id      0
order_amount  0
total_items   0
payment_method 0
created_at    0
dtype: int64
```

```
[9]: len(d1) # 5000 rows
```

```
[9]: 5000
```

```
[10]: d1.mean() # all the means
```

```
C:\Users\prate\anaconda3\envs\shopify\lib\site-packages\ipykernel_launcher.py:1:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with
'numeric_only=None') is deprecated; in a future version this will raise
TypeError. Select only valid columns before calling the reduction.
    """Entry point for launching an IPython kernel.
```

```
[10]: order_id      2500.5000
shop_id      50.0788
user_id      849.0924
order_amount  3145.1280
total_items   8.7872
```

```
dtype: float64
```

Notice the average of order_amount. it is 3145.1280. The average of order_amount was considered as AOV (average order value). This is not considered an accurate calculator of AOV.

A more accurate approach for $AOV = \text{sum of order_amount} / \text{sum of total_items}$

There are a couple of approaches that can be considered. * 1 - We could calculate the total order_amount and divide it by the sum of total_items. * 2 - We could also calculate the individual averages(create a new column avg_aovs = order_amount/total_amount) and then further take the average of this new column. When we say 'averages' we could do 2 things: — 2a - mean of the avg_aovs — 2b - median of the avg_aovs

1.3 Analysis and Observation

```
[11]: AOV_1 = d1['order_amount'].sum() / d1['total_items'].sum()  
AOV_1
```

```
[11]: 357.92152221412965
```

this is without taking outliers into consideration. Let us look for outliers in the order_amount and total_items features.

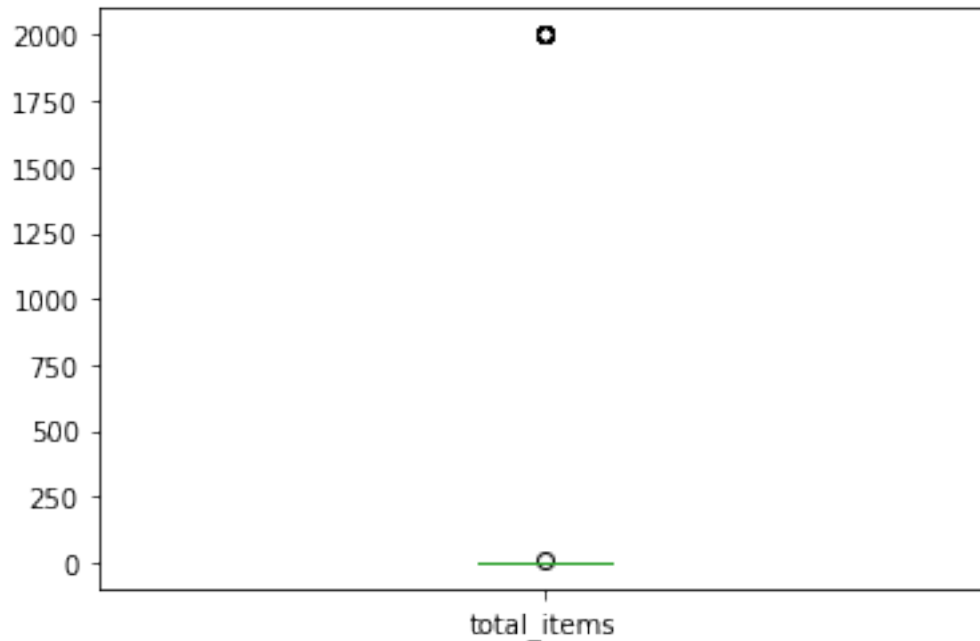
```
[12]: # boxplots for order_amount and total_items  
d1['order_amount'].plot(kind = 'box')
```

```
[12]: <AxesSubplot:>
```



```
[13]: d1['total_items'].plot(kind = 'box')
```

```
[13]: <AxesSubplot:>
```



The boxplots suggests presence of outliers. Its not a great practice to simply remove outliers as they can potentially showcase interesting insights of the situation. It is a good practice to always first investigate the outliers and then make a more informed decision as to how to preceed further. For the sake of this project, since we have limited information, let us attempt to remove the outliers and see.

We define a function ‘outliers’ which we will use to create a list of indcies that contain outliers for the respective feature. We are developing a function since we might need to perform this a number of times. this function will return a list of indicies that contain the outlier for respective feature.

In addition we develop another function, ‘remove’. this function, will take returned list from ‘outliers’ function and return a clean dataframe without the outliers.

Observations that are significantly away from the rest of the data are called outliers. Generally, values beyond 3 standard deviations are considered outliet. We will follow this rule of thumb for this project.

```
[17]: def outliers(df, ft):  
      Q1 = df[ft].quantile(0.25)    # defining the 1st quantile  
      Q3 = df[ft].quantile(0.75)    # defining the 3rd quantile  
      IQR = Q3 - Q1  
      lower_bound = Q1 - 1.5 * IQR  # 1.5 + 1.5 = 3 standard deviation  
      upper_bound = Q3 + 1.5 * IQR
```

```

# getting the indexes          # OR operator
ls = df.index[(df[ft] < lower_bound) | (df[ft] > upper_bound)]
# anything lower than the lower bound OR anything greater than the upper
↳ bound

return ls # list

def remove(df, ls):          # ls -> index list
    ls = sorted(set(ls)) # to get unique indices in ascending order.
    df2 = df.drop(ls)      # drop the respective indices, df2 is the clean
↳ dataframe.
    return df2

```

```

[19]: ind_ls_1 = []
      for i in ['order_amount', 'total_items']:
          ind_ls_1.extend(outliers(d1, i))
      ind_ls_1 # these are the indices of the outliers
      d1_clean = remove(d1, ind_ls_1)
      d1_clean

```

```

[19]:
      order_id  shop_id  user_id  order_amount  total_items  payment_method \
0           1         53      746           224           2           cash
1           2         92      925            90           1           cash
2           3         44      861           144           1           cash
3           4         18      935           156           1  credit_card
4           5         18      883           156           1  credit_card
...         ...         ...         ...         ...         ...
4995        4996         73      993           330           2           debit
4996        4997         48      789           234           2           cash
4997        4998         56      867           351           3           cash
4998        4999         60      825           354           2  credit_card
4999        5000         44      734           288           2           debit

      created_at
0   2017-03-13 12:36:56
1   2017-03-03 17:38:52
2   2017-03-14 4:23:56
3   2017-03-26 12:43:37
4   2017-03-01 4:35:11
...         ...
4995  2017-03-30 13:47:17
4996  2017-03-16 20:36:16
4997  2017-03-19 5:42:42
4998  2017-03-16 14:51:18
4999  2017-03-18 15:48:18

```

[4859 rows x 7 columns]

```
[21]: aov2 = d1_clean['order_amount'].sum() / d1_clean['total_items'].sum()
aov2
```

[21]: 150.60816800337696

```
[22]: aov3 = d1_clean['order_amount'].mean() / d1_clean['total_items'].mean()
aov3
```

[22]: 150.60816800337696

```
[24]: # aov4 = d1_clean['order_amount'].mode() / d1_clean['total_items'].mode()
# aov4
```

[24]: 0 76.5
dtype: float64

Another approach is to calculate respective averages and then calculate the average of the average order value. for this we need to create a new column which is the ratio of order_amount and total_items. We will name the new column 'avg'.

it is generally preferred that the original imported dataframe should not be manipulated. Since we are attempting to create a new column avg_aovs, let us first create a copy of the original dataframe and work on that one.

```
[27]: d2 = d1.copy()
d2.head(2)
```

```
[27]:   order_id  shop_id  user_id  order_amount  total_items  payment_method  \
0         1        53      746          224           2             cash
1         2        92      925           90           1             cash

      created_at
0  2017-03-13 12:36:56
1  2017-03-03 17:38:52
```


```
[28]: d2['avg'] = d2['order_amount'] / d2['total_items']
d2.head(2)
```

```
[28]:   order_id  shop_id  user_id  order_amount  total_items  payment_method  \
0         1        53      746          224           2             cash
1         2        92      925           90           1             cash

      created_at    avg
0  2017-03-13 12:36:56  112.0
1  2017-03-03 17:38:52   90.0
```

```
[30]: # let us look for outliers of avg column. and attempt to remove them
d2['avg'].plot(kind = 'box')
```

```
[30]: <AxesSubplot:>
```



```
[31]: # removing outliers:
ind_ls_2 = []
ind_ls_2.extend(outliers(d2, 'avg'))
ind_ls_2 # these are the indices of the outliers
d2_clean = remove(d2, ind_ls_2)
d2_clean
```

```
[31]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	\
0	1	53	746	224	2	cash	
1	2	92	925	90	1	cash	
2	3	44	861	144	1	cash	
3	4	18	935	156	1	credit_card	
4	5	18	883	156	1	credit_card	
...	
4995	4996	73	993	330	2	debit	
4996	4997	48	789	234	2	cash	
4997	4998	56	867	351	3	cash	
4998	4999	60	825	354	2	credit_card	
4999	5000	44	734	288	2	debit	

```

created_at    avg

```

```


0      2017-03-13 12:36:56 112.0
1      2017-03-03 17:38:52  90.0
2      2017-03-14  4:23:56 144.0
3      2017-03-26 12:43:37 156.0
4      2017-03-01  4:35:11 156.0
...
4995   2017-03-30 13:47:17 165.0
4996   2017-03-16 20:36:16 117.0
4997   2017-03-19  5:42:42 117.0
4998   2017-03-16 14:51:18 177.0
4999   2017-03-18 15:48:18 144.0

```

[4903 rows x 8 columns]

```
[32]: d2_clean['avg'].plot(kind = 'box')
```

[32]: <AxesSubplot:>



```
[34]: aov5 = d2_clean['avg'].mean()
aov5
```

[34]: 150.40016316540894

```
[36]: aov6 = d2_clean['avg'].median()
aov6
```



```
[36]: 153.0
```

```
[37]: # aov7 = d2_clean['avg'].mode()  
# aov7
```

```
[37]: 0    153.0  
dtype: float64
```

1.4 Conclusion:

We are only considering mean and median since the feature is quantitative. Had the feature been qualitative, we would have considered mode.

Based on our analysis, \$150 dollar seem to be a more accurate Average of value for the given situation.

2 Question 2

2.0.1 Part A: How many orders were shipped by speedy express in total?

Orders table contain details of orders. This is our main table. We need to count all the orders shipped by 'Speedy Express'. The table consists of ShipperID. We need to identify which Shipper ID represents Speedy Express. That detail is in the 'Shipper' table. We join the From 'Shippers' table, we get that ShipperID = 1 is for Speedy Express.

```
[42]: '''  
SELECT ShipperName , COUNT(*)  
FROM [Orders]  
JOIN Shippers ON Orders.ShipperID = Shippers.ShipperID  
WHERE ShipperName = 'Speedy Express';  
'''
```

```
[42]: "\nSELECT ShipperName , COUNT(*)\n      FROM [Orders]\n      JOIN Shippers ON  
Orders.ShipperID = Shippers.ShipperID\n      WHERE ShipperName = 'Speedy  
Express';\n"
```

ANSWER: Speedy Express shipped a total of **54 orders**.

2.0.2 Part B: What is the last name of the employee with the most orders?

This time we need to extract the last name of the employee. This information is available in 'Employees' table. The table also contains the 'EmployeeID' of the respective employee. We will match (join) it with the EmployeeID from 'Orders' table.

```
[44]: '''  
SELECT LastName, COUNT(LastName)  
FROM [Orders]  
JOIN Employees ON Orders.EmployeeID = Employees.EmployeeID  
GROUP BY LastName
```

```
ORDER BY COUNT(LastName) DESC
LIMIT 1
```

```
'''
```

```
[44]: '\nSELECT LastName, COUNT(LastName)\n      FROM [Orders]\n      JOIN Employees ON\nOrders.EmployeeID = Employees.EmployeeID\n      GROUP BY LastName\n      ORDER BY\nCOUNT(LastName) DESC\n      LIMIT 1\n      \n'
```

ANSWER: The last name of the employee with the most orders is **‘Peacock’**

2.0.3 Part C: What product was ordered the most by customers in Germany?

For this questions, we need to merge a number of tables since different tables contain different interconnected information. In total, we used Orders, OrderDetails, Products and Customers tables. We selected only Germany. Since we needs, maximum orders per country, we will group the ProductNames and add the quantity. To get the maximum Orders, we will sort the sum in descending order and extract only the 1st row.

```
[46]: '''
```

```
SELECT ProductName, SUM(Quantity)
FROM [Orders]
JOIN OrderDetails ON Orders.OrderID = OrderDetails.OrderID
JOIN Products ON Products.ProductID = OrderDetails.ProductID
JOIN Customers ON Orders.CustomerID = Customers.CustomerID
WHERE Country = 'Germany'
GROUP BY ProductName
ORDER BY SUM(Quantity) DESC
LIMIT 1
'''
```

```
[46]: "\nSELECT ProductName, SUM(Quantity)\n      FROM [Orders]\n      JOIN OrderDetails\nON Orders.OrderID = OrderDetails.OrderID\n      JOIN Products ON\nProducts.ProductID = OrderDetails.ProductID\n      JOIN Customers ON\nOrders.CustomerID = Customers.CustomerID\n      WHERE Country = 'Germany'\n      GROUP BY ProductName\n      ORDER BY SUM(Quantity) DESC\n      LIMIT 1\n      "
```

ANSWER: **Boston Crab Meat** was ordered the most by customers in Germany with a grand total of 160 orders.

```
[ ]:
```